


Introduction to Deep Learning

Bazyli Polednia 2020

Plan for today

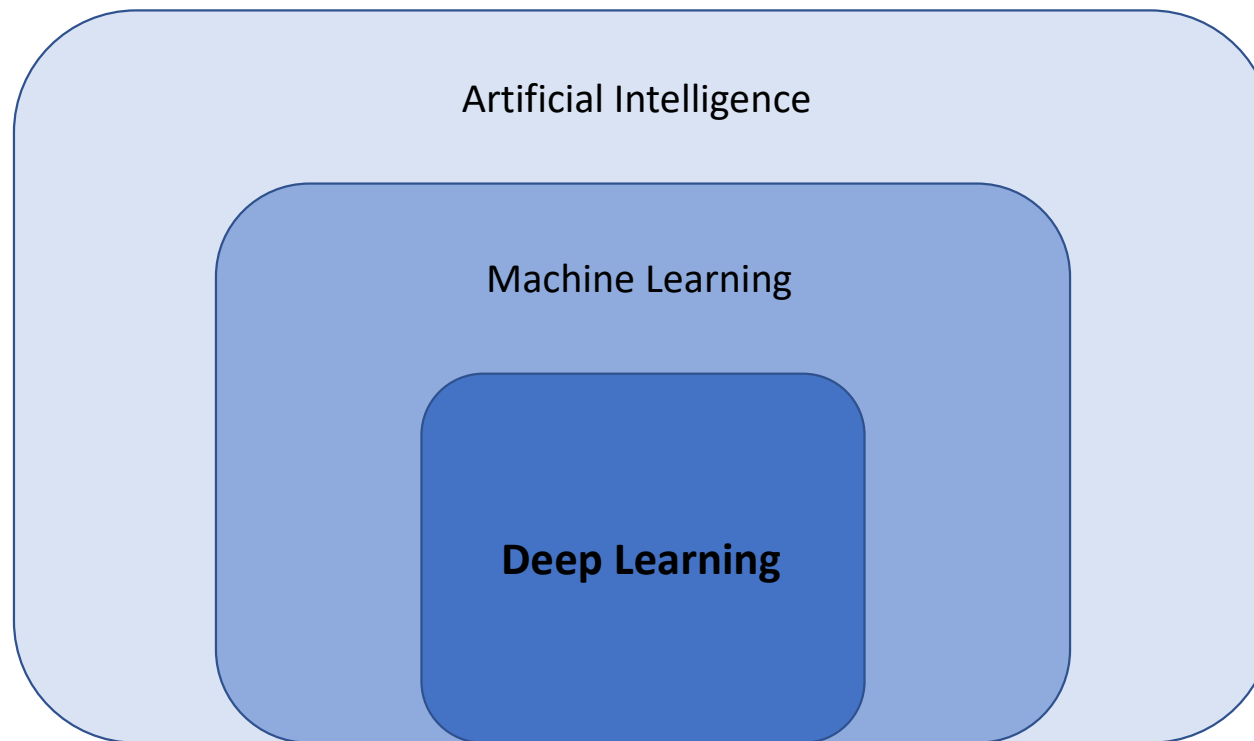
Total: 1h



Resources


- [MIT 6.S191 – Introduction to Deep Learning](#)
- [„Deep Learning with Python” – Francois Chollet](#)
- [Stanford CS230 – Deep Learning](#)
- [Tensorflow Tutorials](#)
- [Kaggle Competitions](#)

Deep Learning – what it actually is?

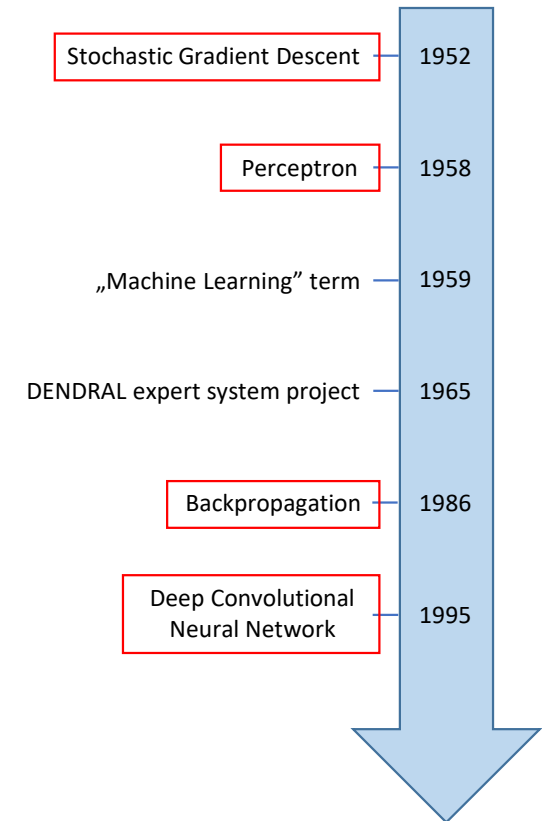



Artificial Intelligence – uses various algorithms to mimic human behaviour. Usually consists of handcrafted rules – such paradigm is called symbolic AI

Machine Learning – instead of following explicit rules provided by programmer, programs learn them by analysing already existing answers to given data



Deep Learning – branch of machine learning using neural networks to recognise patterns in data



Deep Learning Boom in recent years

1. Hardware and software

- Introduction of TPUs and more powerful GPUs
- Parallelization of computations
- Programming interfaces e.g. Nvidia CUDA
- Deep Learning libraries and toolkits e.g. Tensorflow, PyTorch, Keras



2. Datasets and benchmarks


- „Big Data”
- Exponential growth of hardware storage
- Public datasets e.g. ImageNet, Kaggle competitions

IMAGENET

kaggle™

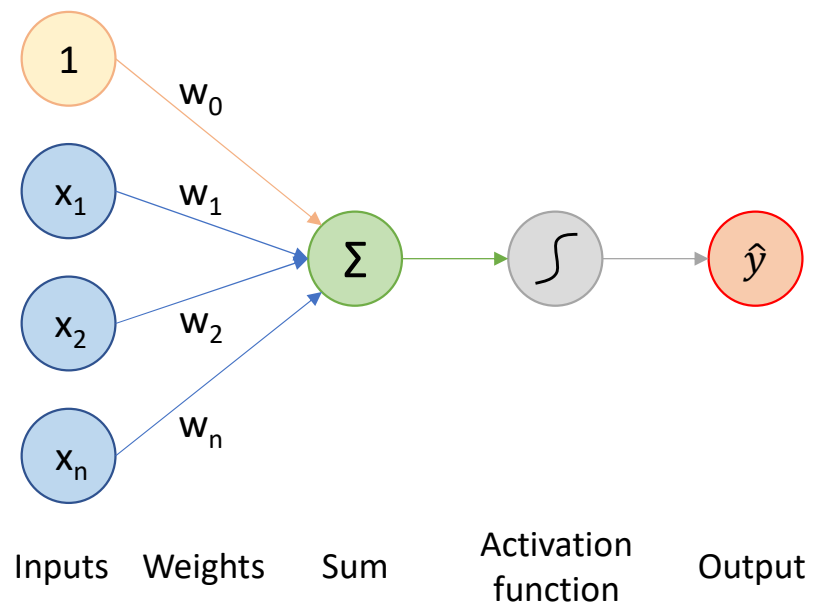
3. Algorithmic advances

- New activation functions
- New optimizers
- Batch optimization
- ...



Perceptron

How is a perceptron built?



Output Activation function Bias Linear combination of inputs

$$\hat{y} = g \left(w_0 + \sum_{j=1}^n x_j w_j \right)$$

$$X = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \quad W = \begin{bmatrix} w_1 \\ \vdots \\ w_n \end{bmatrix}$$

$$\hat{y} = g(w_0 + X^T W)$$

How is a perceptron built – activation function

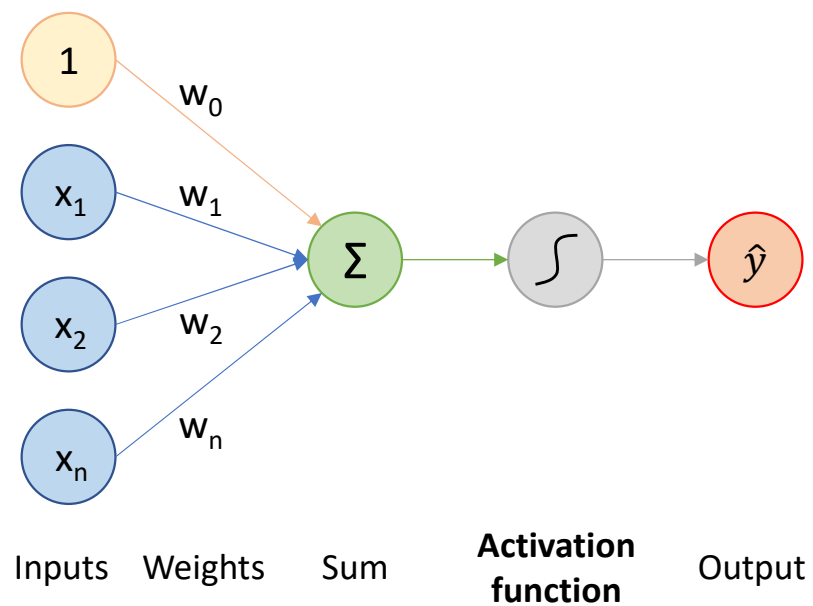


Diagram illustrating the mathematical representation of the perceptron output:

$$\hat{y} = g(w_0 + X^T W)$$


Labels and arrows pointing to the equation:

- Output** (red arrow) points to \hat{y} .
- Activation function** (gray arrow) points to g .
- Bias** (yellow arrow) points to w_0 .
- Linear combination of inputs** (green arrow) points to $X^T W$.


Activation function – what is it?

- Non-linear function
- Transforms linear operations (dot product and addition) of weighted inputs and bias
- Introduces non-linearity to the output


Common examples



ReLU (Rectified Linear Unit)




Sigmoid




Hyperbolic tangent

Activation function – why is it used?



Linearity





Non-linearity

Without non-linear activation function, the perceptron could only learn linear transformations and would only produce linear outputs – which almost never exist in real life


Neural network

Simplified perceptron notation



$$z = w_0 + \sum_{j=1}^n x_j w_j$$

Let's connect two perceptrons to the same inputs



Dense layer

$$z_i = w_{0,i} + \sum_{j=1}^n x_j w_{j,i}$$

Bias


Linear combination of inputs

$$\hat{y}_i = g(z_i)$$

Output

Activation function


Let's connect two perceptrons to the same inputs




$$z_1 = w_{0,1} + \sum_{j=1}^n x_j w_{j,1}$$

$$\hat{y}_1 = g(z_1)$$

Single Layer Neural Network

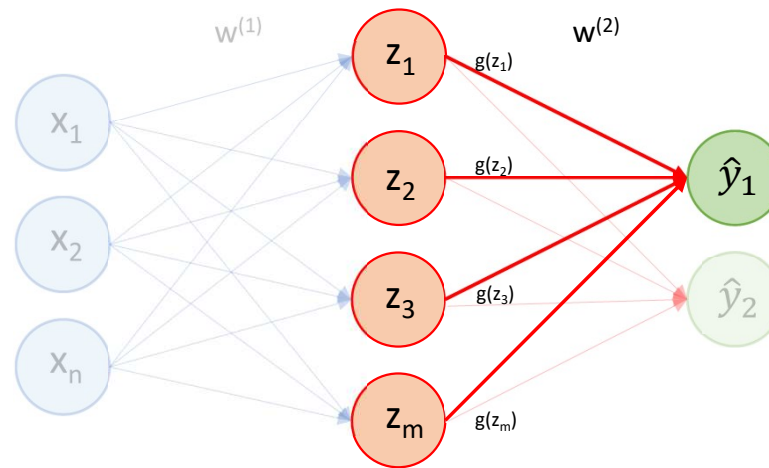


Single Layer Neural Network



$$z_1 = w_{0,1}^{(1)} + \sum_{j=1}^n x_j w_{j,1}^{(1)}$$



Single Layer Neural Network



$$\hat{y}_1 = g \left(w_{0,1}^{(2)} + \sum_{j=1}^m g(z_j) w_{j,1}^{(2)} \right)$$



Example problem

Does the patient have a heart disease?



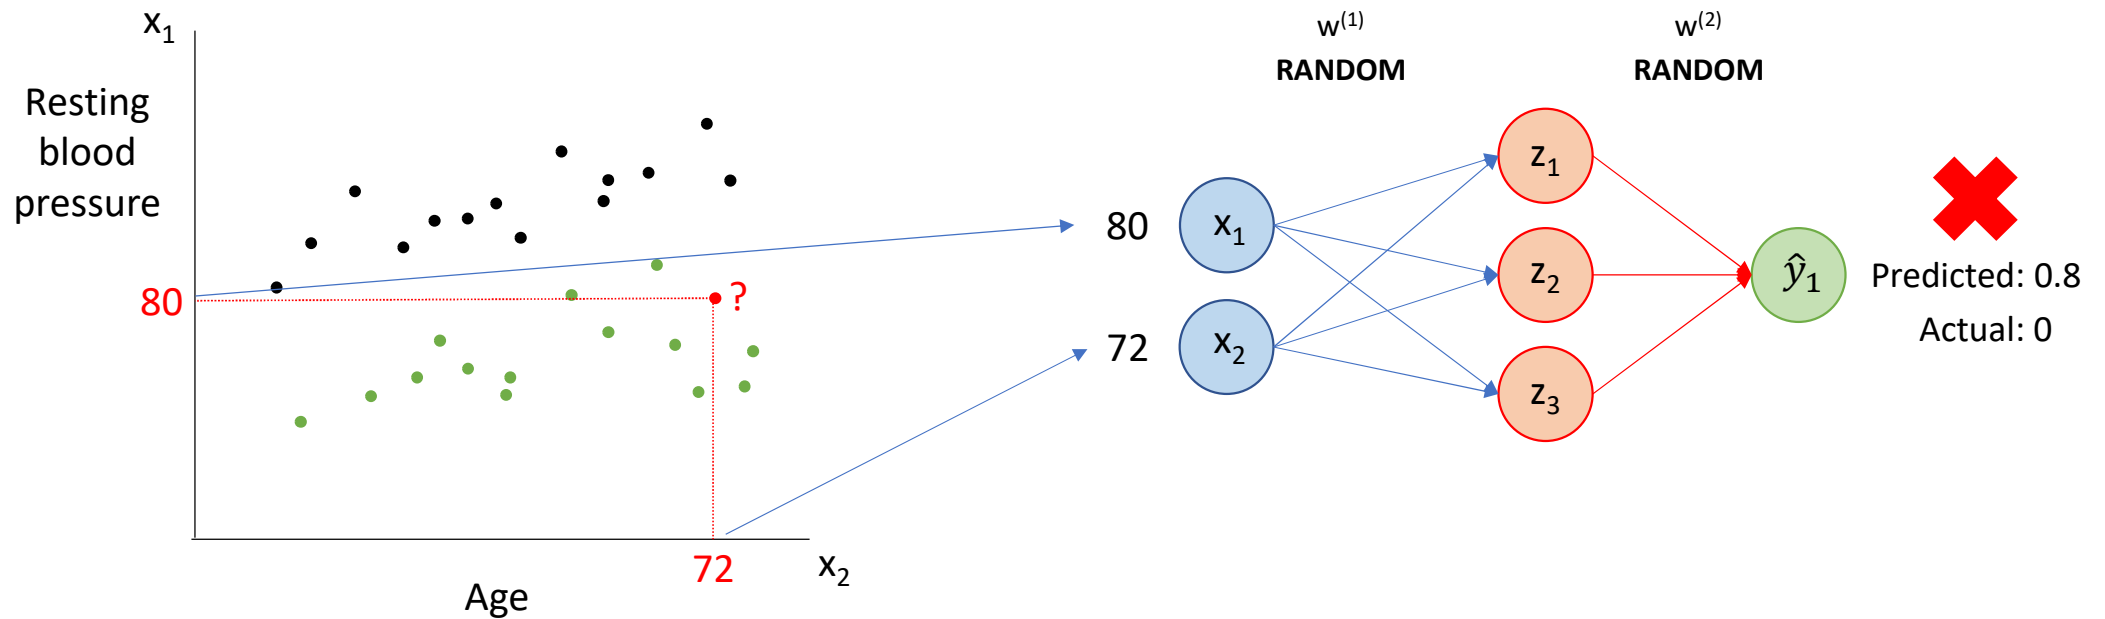
Example problem

Does the patient have a heart disease?




Example problem

Does the patient have a heart disease?



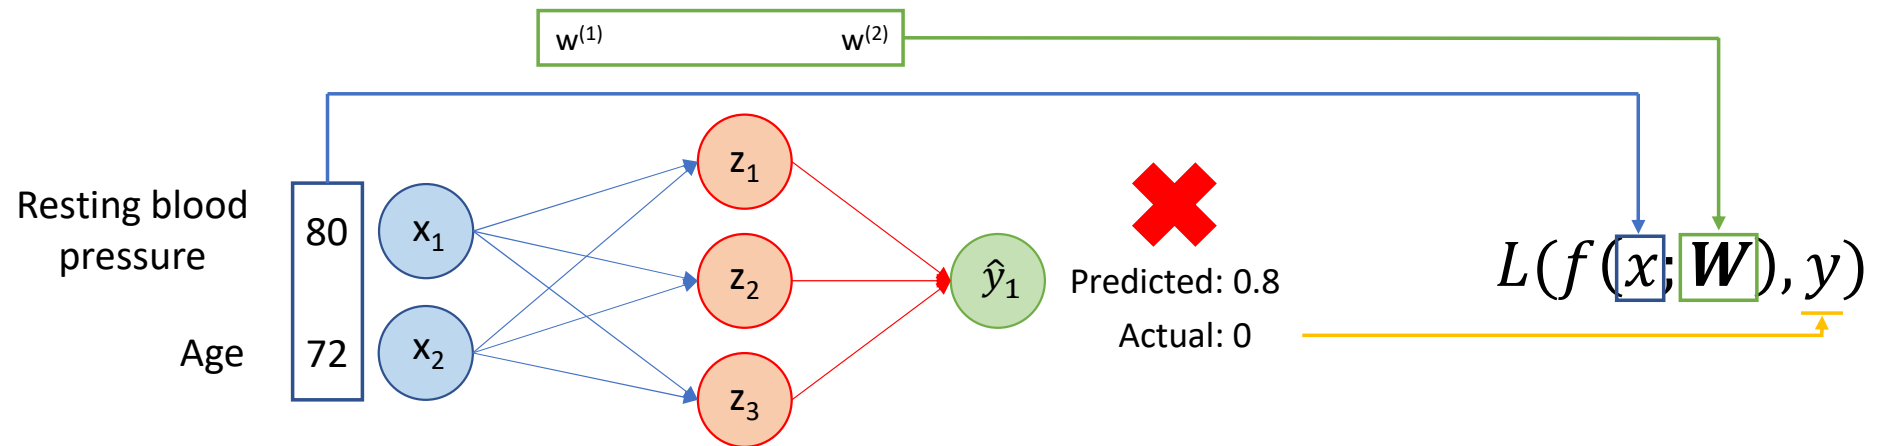
Training the network

Loss function



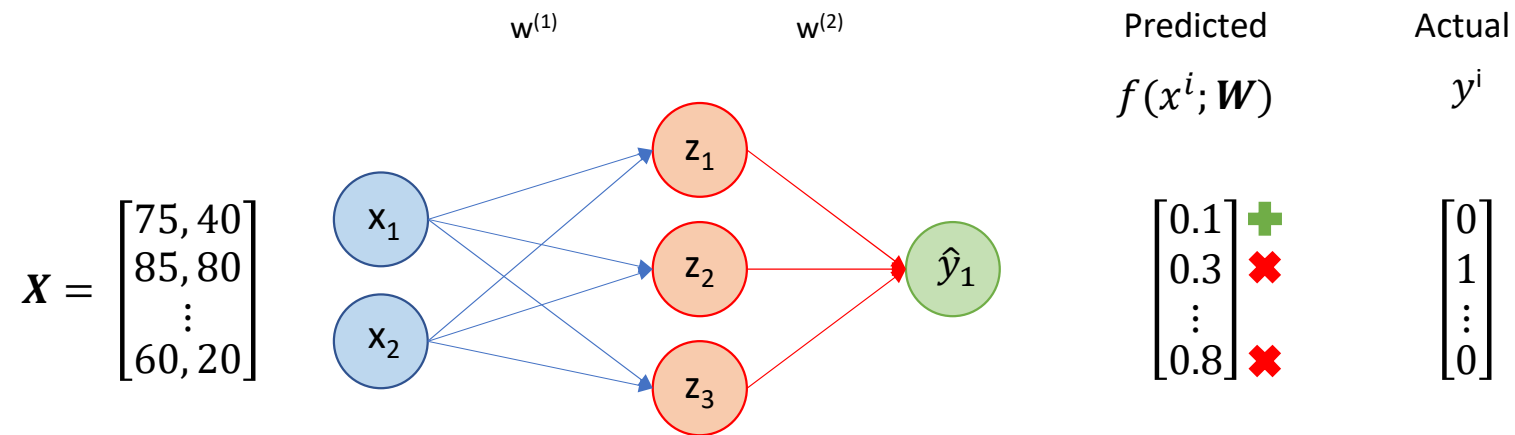
Loss function measures how much our predictions differ from actual results

Loss function



Loss function measures how much our predictions differ from actual results

Empirical loss




$$J(\mathbf{W}) = \frac{1}{n} \sum_{i=1}^n L(f(x^i; \mathbf{W}), y^i)$$


Empirical loss measures loss over the whole dataset, calculating the average of losses for each input

Gradient-based optimization


Problem: minimize loss over weight w




Gradient-based optimization




Gradient-based optimization




Gradient-based optimization



Gradient-based optimization



Gradient-based optimization



Gradient Descent





Image source: „On Why Gradient Descent is Even Needed” Daniel Burkhardt Derigo

Gradient Descent

1. Initialize random weights
2. Loop until convergence:
 - Compute loss function gradient: $\frac{\partial J(W)}{\partial W}$
 - Update weights based on gradient: $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$ η – learning rate

Backpropagation

Problem: compute loss function gradient $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$




$$\frac{\partial J(\mathbf{W})}{\partial w^{(1)}} \quad ? \quad \frac{\partial J(\mathbf{W})}{\partial w^{(2)}}$$

Chain rule

Backpropagation


Problem: compute loss function gradient $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$



$$\frac{\partial J(\mathbf{W})}{\partial w^{(2)}} = ?$$

Backpropagation


Problem: compute loss function gradient $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$



$$\frac{\partial J(\mathbf{W})}{\partial w^{(2)}} = \frac{\partial J(\mathbf{W})}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w^{(2)}}$$

Backpropagation

Problem: compute loss function gradient $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$

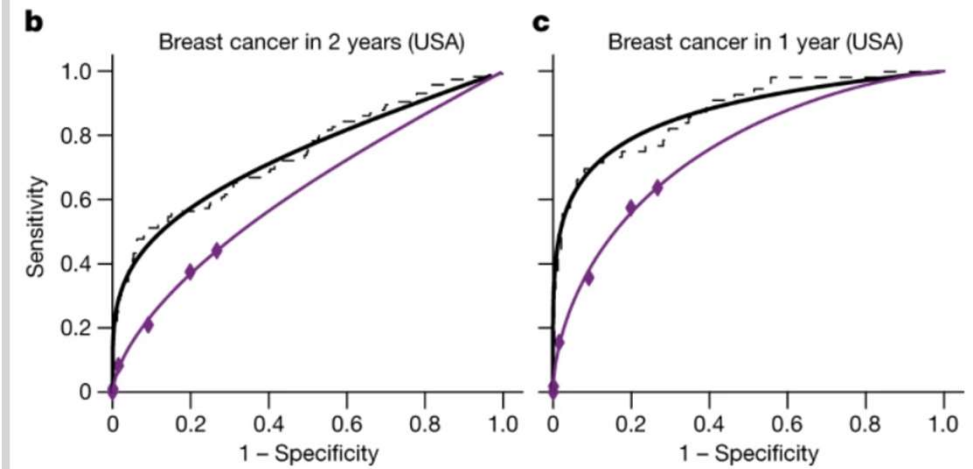



$$\frac{\partial J(\mathbf{W})}{\partial w^{(2)}} = \frac{\partial J(\mathbf{W})}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w^{(2)}}$$

$$\frac{\partial J(\mathbf{W})}{\partial w^{(1)}} = \frac{\partial J(\mathbf{W})}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z} * \frac{\partial z}{\partial w^{(1)}}$$

What can we use
Deep Learning for?

Clinical diagnosis



„International evaluation of an AI system for breast cancer screening” S. M. McKinney, M. Sieniek, V. Godbole, J. Godwin 2020

Real-time object detection




Image source: Detect-Me

Image alteration



Image alteration



And much more...

Questions?

Thank you for attention